

1. Develop simple classification algorithms using a simple 3-layer Neural network, two-layer CNN, and a two-layer LSTM on the following databases: (60 Marks)

a. EHR – Pancreatic

CODE-

### #importing libraries

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import Dense
import numpy as np
from sklearn.impute import SimpleImputer
```

```
# Load the dataset
df = pd.read_csv('EHR_Data.csv')
#print(df)
```

### #Cleaning and Encoding the Data

```
# Binary encode the 'Sensitization type' column
df['Sensitization type'] = df['Sensitization type'].map({'Widespread sensitiza
tion': 1, 'Segmental sensitization': 0})
```

```
# Print the updated dataframe
print(df)
```

```
# replace missing values in 'Gender' column with 'Unknown'
df['Gender'].fillna('Unknown', inplace=True)
```

```
# remove leading/trailing white space and convert to lowercase
df['Gender'] = df['Gender'].str.strip().str.lower()
```

```
# one-hot encode the 'Gender' column
onehot = pd.get_dummies(df['Gender'], prefix='Gender')
```

```
# concatenate the one-hot encoded 'Gender' columns with the original dataset
df = pd.concat([df, onehot], axis=1)
```

```
# drop the original 'Gender' column
df.drop('Gender', axis=1, inplace=True)
```

```
# print the updated dataset
#print(df)
```

```
# replace missing values in 'Etiology' column with 'Unknown'
df['Etiology'].fillna('Unknown', inplace=True)
```

```
# remove leading/trailing white space and convert to lowercase
df['Etiology'] = df['Etiology'].str.strip().str.lower()
```

```

# one-hot encode the 'Etiology' column
onehot = pd.get_dummies(df['Etiology'], prefix='Etiology')

# concatenate the one-hot encoded 'Etiology' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Etiology' column
df.drop('Etiology', axis=1, inplace=True)

# print the updated dataset
#print(df)

# replace missing values in 'Painless/painfull' column with 'Unknown'
df['Painless/painfull'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Painless/painfull'] = df['Painless/painfull'].str.strip().str.lower()

# one-hot encode the 'Painless/painfull' column
onehot = pd.get_dummies(df['Painless/painfull'], prefix='Painless/painfull')

# concatenate the one-
hot encoded 'Painless/painfull' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Painless/painfull' column
df.drop('Painless/painfull', axis=1, inplace=True)

# print the updated dataset
#print(df)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['Duration betwn sx and Dx (months)'] = imputer.fit_transform(df[['Duration
betwn sx and Dx (months)']])

# print the updated dataset
#print(df)

# replace missing values in 'Diabetes' column with 'Unknown'
df['Diabetes'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Diabetes'] = df['Diabetes'].str.strip().str.lower()

# one-hot encode the 'Diabetes' column
onehot = pd.get_dummies(df['Diabetes'], prefix='Diabetes')

# concatenate the one-hot encoded 'Diabetes' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

```

```

# drop the original 'Diabetes' column
df.drop('Diabetes', axis=1, inplace=True)

# print the updated dataset
#print(df)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['Duration betwn CP Dx and diabetes Dx (months)'] = imputer.fit_transform(df[['Duration betwn CP Dx and diabetes Dx (months)']])

# print the updated dataset
#print(df)

# replace missing values in 'Overall_Interventions_yes_no' column with 'Unknown'
df['Overall_Interventions_yes_no'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Overall_Interventions_yes_no'] = df['Overall_Interventions_yes_no'].str.strip().str.lower()

# one-hot encode the 'Overall_Interventions_yes_no' column
onehot = pd.get_dummies(df['Overall_Interventions_yes_no'], prefix='Overall_Interventions_yes_no')

# concatenate the one-hot encoded 'Overall_Interventions_yes_no' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Overall_Interventions_yes_no' column
df.drop('Overall_Interventions_yes_no', axis=1, inplace=True)

# print the updated dataset
#print(df)

# replace missing values in 'Pancreatic_Atrophy' column with 'Unknown'
df['Pancreatic_Atrophy'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Pancreatic_Atrophy'] = df['Pancreatic_Atrophy'].str.strip().str.lower()

# one-hot encode the 'Pancreatic_Atrophy' column
onehot = pd.get_dummies(df['Pancreatic_Atrophy'], prefix='Pancreatic_Atrophy')

# concatenate the one-hot encoded 'Pancreatic_Atrophy' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

```

```

# drop the original 'Pancreatic_Atrophy' column
df.drop('Pancreatic_Atrophy', axis=1, inplace=True)

# replace missing values in 'Dilatation_of_MPD' column with 'Unknown'
df['Dilatation_of_MPD'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Dilatation_of_MPD'] = df['Dilatation_of_MPD'].str.strip().str.lower()

# one-hot encode the 'Dilatation_of_MPD' column
onehot = pd.get_dummies(df['Dilatation_of_MPD'], prefix='Dilatation_of_MPD')

# concatenate the one-
hot encoded 'Dilatation_of_MPD' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Dilatation_of_MPD' column
df.drop('Dilatation_of_MPD', axis=1, inplace=True)

# replace missing values in 'MPD_stricture' column with 'Unknown'
df['MPD_stricture'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['MPD_stricture'] = df['MPD_stricture'].str.strip().str.lower()

# one-hot encode the 'MPD_stricture' column
onehot = pd.get_dummies(df['MPD_stricture'], prefix='MPD_stricture')

# concatenate the one-
hot encoded 'MPD_stricture' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'MPD_stricture' column
df.drop('MPD_stricture', axis=1, inplace=True)

# replace missing values in 'MPD_calculus' column with 'Unknown'
df['MPD_calculus'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['MPD_calculus'] = df['MPD_calculus'].str.strip().str.lower()

# one-hot encode the 'MPD_calculus' column
onehot = pd.get_dummies(df['MPD_calculus'], prefix='MPD_calculus')

# concatenate the one-
hot encoded 'MPD_calculus' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'MPD_calculus' column
df.drop('MPD_calculus', axis=1, inplace=True)

# replace missing values in 'Parenchymal_calcification_' column with 'Unknown'

```

```

df['Parenchymal_calcification_'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Parenchymal_calcification_'] = df['Parenchymal_calcification_'].str.strip().str.lower()

# one-hot encode the 'Parenchymal_calcification_' column
onehot = pd.get_dummies(df['Parenchymal_calcification_'], prefix='Parenchymal_calcification_')

# concatenate the one-hot encoded 'Parenchymal_calcification_' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Parenchymal_calcification_' column
df.drop('Parenchymal_calcification_', axis=1, inplace=True)

# replace missing values in 'Radiation_of_pain' column with 'Unknown'
df['Radiation_of_pain'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Radiation_of_pain'] = df['Radiation_of_pain'].str.strip().str.lower()

# one-hot encode the 'Radiation_of_pain' column
onehot = pd.get_dummies(df['Radiation_of_pain'], prefix='Radiation_of_pain')

# concatenate the one-hot encoded 'Radiation_of_pain' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Radiation_of_pain' column
df.drop('Radiation_of_pain', axis=1, inplace=True)

# replace missing values in 'Development_of_new_areas_of_pain' column with 'Unknown'
df['Development_of_new_areas_of_pain'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Development_of_new_areas_of_pain'] = df['Development_of_new_areas_of_pain'].str.strip().str.lower()

# one-hot encode the 'Development_of_new_areas_of_pain' column
onehot = pd.get_dummies(df['Development_of_new_areas_of_pain'], prefix='Development_of_new_areas_of_pain')

# concatenate the one-hot encoded 'Development_of_new_areas_of_pain' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Development_of_new_areas_of_pain' column
df.drop('Development_of_new_areas_of_pain', axis=1, inplace=True)

```

```

# replace missing values in 'Overall_Pain_Pattern' column with 'Unknown'
df['Overall_Pain_Pattern'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Overall_Pain_Pattern'] = df['Overall_Pain_Pattern'].str.strip().str.lower(
)

# one-hot encode the 'Overall_Pain_Pattern' column
onehot = pd.get_dummies(df['Overall_Pain_Pattern'], prefix='Overall_Pain_Patte
rn')

# concatenate the one-
hot encoded 'Overall_Pain_Pattern' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Overall_Pain_Pattern' column
df.drop('Overall_Pain_Pattern', axis=1, inplace=True)

# replace missing values in 'Pain continuous intermitent' column with 'Unknown'
,
df['Pain continuous intermitent'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Pain continuous intermitent'] = df['Pain continuous intermitent'].str.stri
p().str.lower()

# one-hot encode the 'Pain continuous intermitent' column
onehot = pd.get_dummies(df['Pain continuous intermitent'], prefix='Pain contin
uous intermitent')

# concatenate the one-
hot encoded 'Pain continuous intermitent' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Pain continuous intermitent' column
df.drop('Pain continuous intermitent', axis=1, inplace=True)

# print the updated dataset
#print(df)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['Rumination'] = imputer.fit_transform(df[['Rumination']])

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['Magnification'] = imputer.fit_transform(df[['Magnification']])

```

```

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['Helplessness'] = imputer.fit_transform(df[['Helplessness']])

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['IZBICKI_Final_Score'] = imputer.fit_transform(df[['IZBICKI_Final_Score']])

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['HADS_Depression_Score_'] = imputer.fit_transform(df[['HADS_Depression_Score_']])

# replace missing values in 'HADS_depression_category' column with 'Unknown'
df['HADS_depression_category'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['HADS_depression_category'] = df['HADS_depression_category'].str.strip().str.lower()

# one-hot encode the 'HADS_depression_category' column
onehot = pd.get_dummies(df['HADS_depression_category'], prefix='HADS_depression_category')

# concatenate the one-hot encoded 'HADS_depression_category' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'HADS_depression_category' column
df.drop('HADS_depression_category', axis=1, inplace=True)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['HADS_Anxiety_score'] = imputer.fit_transform(df[['HADS_Anxiety_score']])

# replace missing values in 'HADS_anxiety_category' column with 'Unknown'
df['HADS_anxiety_category'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['HADS_anxiety_category'] = df['HADS_anxiety_category'].str.strip().str.lower()

# one-hot encode the 'HADS_anxiety_category' column

```

```

onehot = pd.get_dummies(df['HADS_anxiety_category'], prefix='HADS_anxiety_category')

# concatenate the one-hot encoded 'HADS_anxiety_category' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'HADS_anxiety_category' column
df.drop('HADS_anxiety_category', axis=1, inplace=True)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['Pain_DETECT_final_score'] = imputer.fit_transform(df[['Pain_DETECT_final_score']])

# replace missing values in 'PainDETECT_typeofpain' column with 'Unknown'
df['PainDETECT_typeofpain'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['PainDETECT_typeofpain'] = df['PainDETECT_typeofpain'].str.strip().str.lower()

# one-hot encode the 'PainDETECT_typeofpain' column
onehot = pd.get_dummies(df['PainDETECT_typeofpain'], prefix='PainDETECT_typeofpain')

# concatenate the one-hot encoded 'PainDETECT_typeofpain' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'PainDETECT_typeofpain' column
df.drop('PainDETECT_typeofpain', axis=1, inplace=True)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['BDI_II_score'] = imputer.fit_transform(df[['BDI_II_score']])

# replace missing values in 'BDI_II_Category' column with 'Unknown'
df['BDI_II_Category'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['BDI_II_Category'] = df['BDI_II_Category'].str.strip().str.lower()

# one-hot encode the 'BDI_II_Category' column
onehot = pd.get_dummies(df['BDI_II_Category'], prefix='BDI_II_Category')

# concatenate the one-hot encoded 'BDI_II_Category' columns with the original dataset

```



```
df = pd.concat([df, onehot], axis=1)

# drop the original 'BDI_II_Category' column
df.drop('BDI_II_Category', axis=1, inplace=True)

# create an instance of SimpleImputer with mean strategy
imputer = SimpleImputer(strategy='mean')

# fill in the missing values in the column with mean value
df['EORTC_Global_Score'] = imputer.fit_transform(df[['EORTC_Global_Score']])

# fill in the missing values in the column with mean value
df['EORTC_Fatigue'] = imputer.fit_transform(df[['EORTC_Fatigue']])

# fill in the missing values in the column with mean value
df['EORTC_Nausea_Vomiting'] = imputer.fit_transform(df[['EORTC_Nausea_Vomiting']])

# fill in the missing values in the column with mean value
df['EORTC_Pain'] = imputer.fit_transform(df[['EORTC_Pain']])

# fill in the missing values in the column with mean value
df['EORTC_Dyspnoea'] = imputer.fit_transform(df[['EORTC_Dyspnoea']])

# fill in the missing values in the column with mean value
df['EORTC_Insomnia'] = imputer.fit_transform(df[['EORTC_Insomnia']])

# fill in the missing values in the column with mean value
df['EORTC_Appetite_loss'] = imputer.fit_transform(df[['EORTC_Appetite_loss']])

# fill in the missing values in the column with mean value
df['EORTC_Constipation'] = imputer.fit_transform(df[['EORTC_Constipation']])

# fill in the missing values in the column with mean value
df['EORTC_Diarrhoea'] = imputer.fit_transform(df[['EORTC_Diarrhoea']])

# fill in the missing values in the column with mean value
df['EORTC_Financial_difficulties'] = imputer.fit_transform(df[['EORTC_Financial_difficulties']])

# fill in the missing values in the column with mean value
df['EORTC_Physical'] = imputer.fit_transform(df[['EORTC_Physical']])

# fill in the missing values in the column with mean value
df['EORTC_Role'] = imputer.fit_transform(df[['EORTC_Role']])

# fill in the missing values in the column with mean value
df['EORTC_Emotional'] = imputer.fit_transform(df[['EORTC_Emotional']])

# fill in the missing values in the column with mean value
df['EORTC_Cognitive'] = imputer.fit_transform(df[['EORTC_Cognitive']])

# fill in the missing values in the column with mean value
df['EORTC_Social'] = imputer.fit_transform(df[['EORTC_Social']])
```

```

# fill in the missing values in the column with mean value
df['PAN28_pancreatic_pain'] = imputer.fit_transform(df[['PAN28_pancreatic_pain']])

# fill in the missing values in the column with mean value
df['PAN28_digestive_sym'] = imputer.fit_transform(df[['PAN28_digestive_sym']])

# fill in the missing values in the column with mean value
df['PAN28_altered_bowel'] = imputer.fit_transform(df[['PAN28_altered_bowel']])

# fill in the missing values in the column with mean value
df['PAN28_Hepatic'] = imputer.fit_transform(df[['PAN28_Hepatic']])

# fill in the missing values in the column with mean value
df['PAN28_Bodyimage'] = imputer.fit_transform(df[['PAN28_Bodyimage']])

# fill in the missing values in the column with mean value
df['PAN28_Activity'] = imputer.fit_transform(df[['PAN28_Activity']])

# fill in the missing values in the column with mean value
df['PAN28_Satisfaction_with_healthcare'] = imputer.fit_transform(df[['PAN28_Satisfaction_with_healthcare']])

# fill in the missing values in the column with mean value
df['PAN28_Sexuality'] = imputer.fit_transform(df[['PAN28_Sexuality']])

# replace missing values in 'P-QST_Yes_No' column with 'Unknown'
df['P-QST_Yes_No'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['P-QST_Yes_No'] = df['P-QST_Yes_No'].str.strip().str.lower()

# one-hot encode the 'P-QST_Yes_No' column
onehot = pd.get_dummies(df['P-QST_Yes_No'], prefix='P-QST_Yes_No')

# concatenate the one-hot encoded 'P-QST_Yes_No' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'P-QST_Yes_No' column
df.drop('P-QST_Yes_No', axis=1, inplace=True)

# replace missing values in 'Sensitization_yes_no' column with 'Unknown'
df['Sensitization_yes_no'].fillna('Unknown', inplace=True)

# remove leading/trailing white space and convert to lowercase
df['Sensitization_yes_no'] = df['Sensitization_yes_no'].str.strip().str.lower()

# one-hot encode the 'Sensitization_yes_no' column
onehot = pd.get_dummies(df['Sensitization_yes_no'], prefix='Sensitization_yes_no')

```

```
# concatenate the one-hot encoded 'Sensitization_yes_no' columns with the original dataset
df = pd.concat([df, onehot], axis=1)

# drop the original 'Sensitization_yes_no' column
df.drop('Sensitization_yes_no', axis=1, inplace=True)

# print the updated dataset
print(df)
```

Output-

	Sensitization type	Duration betwn sx and Dx (months)	\	
0	0	88.000000		
1	0	0.000000		
2	0	0.000000		
3	0	9.000000		
4	0	10.000000		
...	...	...		
369	1	29.004587		
370	1	29.004587		
371	1	29.004587		
372	1	29.004587		
373	1	29.004587		

	Duration betwn CP Dx and diabetes Dx (months)	Rumination	Magnification
0	16.000000	0.0	0.000000
1	0.000000	2.0	0.000000
2	39.902439	1.0	0.000000
3	0.000000	3.0	1.000000
4	39.902439	0.0	0.000000
...	...	...	...
369	6.000000	0.0	3.000000
370	6.000000	11.0	3.867209
371	6.000000	13.0	0.000000
372	46.000000	8.0	6.000000
373	5.000000	13.0	5.000000

	Helplessness	IZBICKI_Final_Score	HADS_Depression_Score_	\	
0	0.0	60.0	0.0		
1	0.0	35.0	0.0		
2	0.0	35.0	0.0		
3	0.0	47.5	2.0		
4	0.0	31.5	0.0		
...	...	...	...		
369	8.0	50.0	5.0		
370	10.0	80.0	9.0		
371	0.0	80.0	9.0		
372	1.0	80.0	9.0		
373	11.0	80.0	9.0		

	HADS_Anxiety_score	Pain_DETECT_final_score	...	\	
0	0.0	13.0	...		
1	0.0	0.0	...		
2	0.0	4.0	...		
3	6.0	5.0	...		
4	0.0	2.0	...		
...	...	...	...		
369	13.0	9.0	...		
370	13.0	7.0	...		

371	13.0	7.0	...
372	13.0	7.0	...
373	13.0	7.0	...

	BDI_II_Category_mild depression	BDI_II_Category_mild mood disturbance	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
..	...	...	
369	0	0	
370	1	0	
371	1	0	
372	1	0	
373	1	0	

	BDI_II_Category_moderate depression	BDI_II_Category_normal	\
0	0	1	
1	0	1	
2	0	1	
3	0	1	
4	0	1	
..	...	...	
369	0	0	
370	0	0	
371	0	0	
372	0	0	
373	0	0	

	BDI_II_Category_severe depression	BDI_II_Category_unknown	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
..	...	...	
369	0	0	
370	0	0	
371	0	0	
372	0	0	
373	0	0	

	P-QST_Yes_No_no	P-QST_Yes_No_unknown	P-QST_Yes_No_yes	\
0	0	1	0	
1	0	1	0	
2	0	1	0	
3	0	1	0	
4	0	1	0	
..	...	...	...	
369	0	0	1	
370	0	0	1	
371	0	0	1	
372	0	0	1	
373	0	0	1	

	Sensitization_yes_no_yes
0	1
1	1
2	1
3	1
4	1
..	...
369	1
370	1

```
371                                     1
372                                     1
373                                     1
```

```
[374 rows x 114 columns]
```

### 3-layer Neural network-

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the model architecture
model = Sequential()
model.add(Dense(units=64, activation='relu', input_dim=X_train.shape[1]))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train the model
history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_test, y_test), callbacks=[early_stopping])

# Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test, y_test)
print('Test loss:', loss)
print('Test accuracy:', accuracy)

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
```

```

plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)

# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()



```

## Output-


```

Epoch 44/100
10/10 [=====] - 0s 7ms/step - loss: 0.0331 - accuracy:
1.0000 - val_loss: 0.6717 - val_accuracy: 0.8800
3/3 [=====] - 0s 4ms/step - loss: 0.6717 - accuracy:
0.8800
Test loss: 0.6716741919517517
Test accuracy: 0.8799999952316284

```



3/3 [=====] - 0s 3ms/step  
Confusion matrix:  
[[ 4 7]  
 [ 2 62]]  
Precision: 0.8985507246376812  
Recall: 0.96875  
F1-score: 0.9323308270676692  
ROC AUC: 0.6661931818181819



## Two-layer CNN -

```
#CNN
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, Conv1D, Flatten, MaxPooling1D
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Reshape input data for CNN
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)
```



```

# Define the model architecture
model = Sequential()
model.add(Conv1D(filters=64, kernel_size=3, activation='relu', input_shape=X_train.shape[1:]))
model.add(MaxPooling1D(pool_size=2))
model.add(Conv1D(filters=32, kernel_size=3, activation='relu'))
model.add(MaxPooling1D(pool_size=2))
model.add(Flatten())
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train the model
history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_data=(X_test, y_test), callbacks=[early_stopping])

# Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test, y_test)
print('Test loss:', loss)
print('Test accuracy:', accuracy)

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)

# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)


```

```
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()
```


Output-

Epoch 31/100  
10/10 [=====] - 0s 18ms/step - loss: 0.0999 - accuracy:  
0.9799 - val\_loss: 0.4460 - val\_accuracy: 0.8667  
3/3 [=====] - 0s 6ms/step - loss: 0.4460 - accuracy:  
0.8667  
Test loss: 0.44604364037513733  
Test accuracy: 0.8666666746139526





3/3 [=====] - 0s 5ms/step  
Confusion matrix:  
[[ 1 10]  
[ 0 64]]  
Precision: 0.8648648648648649  
Recall: 1.0  
F1-score: 0.927536231884058  
ROC AUC: 0.5454545454545454



## A two-layer LSTM -

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, LSTM
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Resample the training data to address class imbalance
X_train_resampled, y_train_resampled = resample(X_train[y_train==0], y_train[y_train==0],
                                                replace=True, n_samples=X_train[y_train==1].shape[0],
                                                random_state=42)
X_train_resampled = np.concatenate([X_train_resampled, X_train[y_train==1]], axis=0)
y_train_resampled = np.concatenate([y_train_resampled, y_train[y_train==1]], axis=0)

# Define the model architecture
model = Sequential()
model.add(LSTM(units=64, return_sequences=True, input_shape=(X_train.shape[1], 1)))
model.add(LSTM(units=32))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train the model
history = model.fit(X_train_resampled.reshape(-1, X_train.shape[1], 1), y_train_resampled, epochs=100, batch_size=32, validation_data=(X_test.reshape(-1, X_train.shape[1], 1), y_test), callbacks=[early_stopping])
```

```

# Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test.reshape(-1, X_train.shape[1], 1), y_test)
print('Test loss:', loss)
print('Test accuracy:', accuracy)

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test.reshape(-1, X_train.shape[1], 1))
y_pred = (y_pred > 0.5)



# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()


```

# Output-

Epoch 14/100  
17/17 [=====] - 2s 103ms/step - loss: 0.4225 - accuracy: 0.8127 - val\_loss: 0.5832 - val\_accuracy: 0.6267  
3/3 [=====] - 0s 26ms/step - loss: 0.5832 - accuracy: 0.6267  
Test loss: 0.5832009315490723  
Test accuracy: 0.6266666650772095



Confusion matrix:  
[[11 0]  
[28 36]]  
Precision: 1.0  
Recall: 0.5625  
F1-score: 0.72  
ROC AUC: 0.78125



**2. Data is imbalanced. Use the strategy to overcome the challenge of imbalance  
(a) undersampling, (b) oversampling**

### **(a) Undersampling applied to simple 3-layer Neural network-**

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values
```

```

# Undersample the majority class
from imblearn.under_sampling import RandomUnderSampler
rus = RandomUnderSampler(random_state=42)
X_resampled, y_resampled = rus.fit_resample(X, y)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
test_size=0.2, random_state=42)

# Define the model architecture
model = Sequential()
model.add(Dense(units=64, activation='relu', input_dim=X_train.shape[1]))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy
'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train the model
history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_da
ta=(X_test, y_test), callbacks=[early_stopping])

# Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test, y_test)
print('Test loss:', loss)
print('Test accuracy:', accuracy)

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)

```



```

y_pred = (y_pred > 0.5)

# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()



```

## Output-


```

Epoch 10/100
3/3 [=====] - 0s 31ms/step - loss: 0.9694 - accuracy:
0.6471 - val_loss: 3.4241 - val_accuracy: 0.5000
1/1 [=====] - 0s 44ms/step - loss: 3.4241 - accuracy:
0.5000
Test loss: 3.4240646362304688
Test accuracy: 0.5

```



```
1/1 [=====] - 0s 85ms/step
Confusion matrix:
[[5 2]
 [7 4]]
Precision: 0.6666666666666666
Recall: 0.36363636363636365
F1-score: 0.4705882352941177
ROC AUC: 0.538961038961039
```



## (b) Oversampling applied to simple 3-layer Neural network-

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

# Oversample the minority class
from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(X, y)

# Split the dataset into training and testing sets
```

```

X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
test_size=0.2, random_state=42)

# Define the model architecture
model = Sequential()
model.add(Dense(units=64, activation='relu', input_dim=X_train.shape[1]))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy
'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train the model
history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_da
ta=(X_test, y_test), callbacks=[early_stopping])

# Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test, y_test)
print('Test loss:', loss)
print('Test accuracy:', accuracy)

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)

# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f
1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)


```


```
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()
```

Output-

Epoch 24/100  
17/17 [=====] - 0s 7ms/step - loss: 0.0399 - accuracy: 0.9962 - val\_loss: 0.2188 - val\_accuracy: 0.9248  
5/5 [=====] - 0s 4ms/step - loss: 0.2188 - accuracy: 0.9248  
Test loss: 0.21878783404827118  
Test accuracy: 0.9248120188713074





5/5 [=====] - 0s 4ms/step

Confusion matrix:


```
[[58  0]
 [10 65]]
```

Precision: 1.0

Recall: 0.8666666666666667

F1-score: 0.9285714285714286

ROC AUC: 0.9333333333333333



**3. Try (a) k-fold cross-validation, (b) leave one subject out cross-validation and (c) k-stratified cross-validation.**

### **(a) k-fold cross-validation to simple 3-layer Neural network-**

```
#K-Fold CV
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, KFold
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

# Oversample the minority class
from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(X, y)

# Set up k-fold cross-validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Define the model architecture
model = Sequential()
model.add(Dense(units=64, activation='relu', input_dim=X_resampled.shape[1]))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train and evaluate the model using k-fold cross-validation
acc_scores = []
loss_scores = []
for train_idx, test_idx in kf.split(X_resampled, y_resampled):
    X_train, X_test = X_resampled[train_idx], X_resampled[test_idx]
    y_train, y_test = y_resampled[train_idx], y_resampled[test_idx]
```

```

    history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation
n_data=(X_test, y_test), callbacks=[early_stopping], verbose=1)
    loss, accuracy = model.evaluate(X_test, y_test, verbose=0)
    acc_scores.append(accuracy)
    loss_scores.append(loss)

# Compute and print performance metrics
print('Accuracy:', np.mean(acc_scores))
print('Loss:', np.mean(loss_scores))

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)

# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f
1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')



```



```
plt.show()
```

Output-


Epoch 100/100  
17/17 [=====] - 0s 6ms/step - loss: 1.1074e-04 - accuracy: 1.0000 - val\_loss: 1.0520e-04 - val\_accuracy: 1.0000  
Accuracy: 0.9909774541854859  
Loss: 0.028932575223734602



```

Confusion matrix:
[[72  0]
 [ 0 60]]
Precision: 1.0
Recall: 1.0
F1-score: 1.0
ROC AUC: 1.0

```



## (b) leave one subject out cross-validation to simple 3-layer Neural network-

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import LeaveOneGroupOut
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
from imblearn.over_sampling import RandomOverSampler
from sklearn.utils import shuffle

# Load the dataset
data = df

# Sort the data by subject id
data = data.sort_values('Sensitization type')

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values
groups = data.iloc[:, 1].values

```

```

# Define the model architecture
model = Sequential()
model.add(Dense(units=64, activation='relu', input_dim=X.shape[1]))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Set up leave-one-subject-out cross-validation
logo = LeaveOneGroupOut()

# Train and evaluate the model with leave-one-subject-out cross-validation
for train_index, test_index in logo.split(X, y, groups=groups):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    # Shuffle the training set
    X_train, y_train = shuffle(X_train, y_train, random_state=42)

    # Oversample the minority class in the training set
    ros = RandomOverSampler(random_state=42)
    X_train_resampled, y_train_resampled = ros.fit_resample(X_train, y_train)

    # Train the model
    history = model.fit(X_train_resampled, y_train_resampled, epochs=100, batch_size=32, validation_data=(X_test, y_test), callbacks=[early_stopping])

    # Evaluate the model on the test set
    loss, accuracy = model.evaluate(X_test, y_test)
    print('Test loss:', loss)
    print('Test accuracy:', accuracy)

    # Make predictions on the test set
    y_pred = model.predict(X_test)
    y_pred = (y_pred > 0.5)

    # Compute performance metrics
    from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
    cm = confusion_matrix(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    fpr, tpr, thresholds = roc_curve(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred)

    print('Confusion matrix:\n', cm)
    print('Precision:', precision)

```

```

print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()


```

## Output-

```

Epoch 7/100
16/16 [=====] - 0s 6ms/step - loss: 0.1053 - accuracy:
0.9738 - val_loss: 1.1759 - val_accuracy: 0.8061
4/4 [=====] - 0s 3ms/step - loss: 1.1759 - accuracy:
0.8061
Test loss: 1.1759247779846191
Test accuracy: 0.8061224222183228
4/4 [=====] - 0s 2ms/step
Confusion matrix:
[[ 5 10]
 [ 9 74]]
Precision: 0.8809523809523809
Recall: 0.891566265060241
F1-score: 0.8862275449101796
ROC AUC: 0.6124497991967872

```



### (c) k-stratified cross-validation to simple 3-layer Neural network-

```
#stratified cross-validation
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, StratifiedKFold
from keras.models import Sequential
from keras.layers import Dense
from keras.callbacks import EarlyStopping
from sklearn.utils import resample

# Load the dataset
data = df

# Split the dataset into X (independent variables) and y (dependent variable)
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

# Define the number of splits for stratified cross-validation
n_splits = 5

# Define the model architecture
model = Sequential()
model.add(Dense(units=64, activation='relu', input_dim=X.shape[1]))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Set up early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Set up stratified cross-validation
skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=42)

# Train and evaluate the model with stratified cross-validation
for train_index, test_index in skf.split(X, y):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    # Oversample the minority class in the training set
    ros = RandomOverSampler(random_state=42)
    X_train_resampled, y_train_resampled = ros.fit_resample(X_train, y_train)

    # Train the model
    history = model.fit(X_train_resampled, y_train_resampled, epochs=100, batch_size=32, validation_data=(X_test, y_test), callbacks=[early_stopping])

    # Evaluate the model on the test set
    loss, accuracy = model.evaluate(X_test, y_test)
    print('Test loss:', loss)
```

```

    print('Test accuracy:', accuracy)

# Plot accuracy and loss curves
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)

# Compute performance metrics
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, roc_curve, roc_auc_score
cm = confusion_matrix(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)

print('Confusion matrix:\n', cm)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1)
print('ROC AUC:', roc_auc)
plt.plot(fpr, tpr, 'b', label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc='lower right')
plt.show()

```



## Output-

```

Epoch 7/100
17/17 [=====] - 0s 6ms/step - loss: 0.0091 - accuracy:
1.0000 - val_loss: 0.0179 - val_accuracy: 1.0000
3/3 [=====] - 0s 5ms/step - loss: 0.0179 - accuracy:
1.0000
Test loss: 0.017888842150568962

```

Test accuracy: 1.0



3/3 [=====] - 0s 4ms/step  
Confusion matrix:  
[[ 8 0]  
 [ 0 66]]  
Precision: 1.0  
Recall: 1.0  
F1-score: 1.0  
ROC AUC: 1.0

